

Talk to Data

5-Minute Submission Video — Complete Production Guide

From zero recording experience to a polished, engaging final video.
Step-by-step. Word-for-word script. Editing checklist. Done in one weekend.

Project: NatWest Code for Purpose — Theme 1

Presenters: Friend 1 (Naman) and Friend 2 (your friend)

Target length: 4:50 (under 5:00)

Format: Webcam side-by-side + Screen recording PiP

Section 0 — How to use this guide

This guide walks you and your friend through the entire process of producing a 5-minute submission video for the project — from the moment you set up the room to the moment you click Submit. It is intentionally exhaustive: read it once end-to-end, then keep it open during the shoot.

Roles used everywhere in this guide:

- **Friend 1 (Naman)** — primary narrator and demo driver.
- **Friend 2 (your friend)** — co-host, problem framer, and tech-stack speaker.
- **BOTH** — speaking together (intro hook and outro).

Total time you should plan for:

- Pre-production setup: 1.5 hours
- Recording (with retakes): 1.5–2 hours
- Editing: 2–3 hours
- QC + export + upload: 30 minutes

So roughly half a day if you're focused. Less if you've recorded videos before.

Note: *If you only have one camera, record the two of you sitting side-by-side on a single webcam. If you have two laptops/cameras, use both — the editing tutorial covers both layouts.*

Section 1 — Equipment & Software Checklist

Hardware

- 1 or 2 laptops (one runs Streamlit + screen recording; another can be a second webcam).
- Phone tripod (only if you're using your phone as a second webcam).
- Earphones with inline mic, OR a USB mic — anything is better than the laptop's built-in mic.
- A quiet room — close the door, switch fans/AC off if loud.
- A simple background — plain wall, bookshelf, or curtain. Avoid moving objects behind you.
- A lamp or ring light positioned in front of your face (not behind).

Software (all free)

- **OBS Studio** — recording. Captures screen + webcam(s) + mic in one go. obsproject.com
- **DaVinci Resolve** — editing. Pro-grade, free. blackmagicdesign.com
- **CapCut for Mac** — easier editor if Resolve feels heavy. capcut.com
- **Cursor or VS Code** — to show the code on screen during the demo.
- **Chrome / Brave** — to run the Streamlit app and show slides.
- **Google Slides or Keynote** — for two simple slides (title card + tech stack).

Royalty-free assets (for music and minor SFX)

- Music: pixabay.com/music or YouTube Audio Library — search 'corporate inspiring' or 'soft background'.
- SFX (whoosh / click / pop): pixabay.com/sound-effects — pick gentle, short ones.
- Stock B-roll (optional): mixkit.co or pexels.com/videos.

Phone-as-a-webcam options (if you don't have a second camera)

- Mac: use macOS Continuity Camera (System Settings → General → AirDrop & Handoff).
- Cross-platform: install **Camo** or **iVCam** on phone + companion app on laptop.

Section 2 — Pre-Production Setup (90 minutes)

2.1 Workspace setup

- Clean the desk — only the laptop, mic, and a small water bottle visible.
- Sit so the light source (window or lamp) is in front of you, slightly to one side at 45°.
- Both presenters: solid-coloured shirts. Avoid stripes, logos, and pure white (overexposes).
- Decide framing: head + shoulders, eyes about 1/3 down from the top of the frame.
- If sitting side-by-side, leave a small gap between shoulders so each face is fully visible.

2.2 Computer prep

- Close every app you don't need — Slack, WhatsApp Web, mail, Spotify, etc.
- Mac: System Settings → Focus → enable Do Not Disturb. iPhones too.
- Set browser to Full Screen, hide bookmarks bar (View → Hide Bookmarks Bar).
- Hide desktop icons: Mac — Finder Settings → General → uncheck 'Show items'. Or use a clean Desktop wallpaper.
- Increase Terminal/Cursor font size to 16–18pt so it's readable on video.
- Pick a single light/dark theme — don't switch between them on camera.

2.3 Streamlit and project prep

```
cd /Users/naman/Desktop/Talk-to-Data-main
source .venv/bin/activate
streamlit run src/talk_to_data/app.py
```

- Open <http://localhost:8501> in your browser, full-screen the window.
- Pre-paste your Gemini API key in the right-hand panel and click Save key.
- Have **data/sample_hackathon.csv** ready in Finder (drag-zone visible).
- Have a small sample PDF ready (2–3 pages, any topic) — also visible in Finder.
- Open Cursor (or VS Code) on the project folder. Open these files in this exact tab order:
 - 1. README.md
 - 2. src/talk_to_data/app.py
 - 3. src/talk_to_data/views.py
 - 4. src/talk_to_data/analytics.py
 - 5. src/talk_to_data/llm_layer.py
 - 6. src/talk_to_data/pgvector_store.py

2.4 Slides to prepare (only 2 needed)

- **Slide A — Title card:** 'Talk to Data — Theme 1' with both presenters' names below. Background: deep navy gradient (#0f172a → #1e3a8a).
- **Slide B — Tech stack:** 6 small logo tiles — Python, Streamlit, Pandas, Gemini, Postgres+pgvector, Docker.
- Open both in a separate browser window, presenter view ON, ready to F5 / Cmd+Enter.

2.5 OBS Studio setup

Open OBS, run Auto-Configuration Wizard once, choose **Optimise for recording**. Then create one scene with three sources visible:

1. Scene name: **Main Demo**.
2. Add Source → **Display Capture**. Pick the screen with Streamlit/Cursor.
3. Add Source → **Video Capture Device** for Friend 1's webcam. Resize and place bottom-left.
4. Add Source → **Video Capture Device** for Friend 2's webcam. Resize and place bottom-right.
5. Add Source → **Audio Input Capture**. Pick your USB / earphone mic. Test levels (peaks at -12 to -6 dB).

Optional second scene: **Both Faces** — both webcams large, no screen capture. Use this for intros and outros.

OBS settings to tweak (Settings menu)

- Output → Recording → Format **MP4** (or MKV if you'll remux later), Quality High, encoder Apple VT H.264 (Mac) or NVENC.
- Video → Base canvas 1920×1080, Output 1920×1080, FPS 30.
- Audio → Sample rate 48 kHz, Channels Stereo.
- Hotkeys → set **Cmd+R** for Start/Stop Recording so you don't have to click.

2.6 Practice run

- Read the entire script aloud once — time it. Aim for 4:45–4:55.
- Do a 30-second test recording. Play it back. Check audio is clear and webcam is sharp.
- Adjust lamp / mic / fonts as needed. Re-test.

Section 3 — Role Assignment & Speaking Pattern

The 5 minutes are divided into 9 blocks. Two presenters keep energy alive by handing off, with one moment of dual energy at the start and the very end.

0:00 – 0:15 Hook intro	BOTH (alternating one line each)
0:15 – 0:45 Problem	Friend 2 (sets the stage, slightly serious tone)
0:45 – 1:15 Solution overview	Friend 1 (confident, owns the story)
1:15 – 1:50 Tech stack	Friend 2 (technical credibility)
1:50 – 2:30 Demo: CSV upload + first question	Friend 1 (Friend 2 chips in once)
2:30 – 3:05 Demo: second CSV question	Friend 2 leads, Friend 1 narrates clicks
3:05 – 3:35 Demo: PDF + AI Q&A;	Friend 2 (clean handoff)
3:35 – 4:10 Code walkthrough (brief)	Friend 1 (engineering pride)
4:10 – 4:40 Why this stands out	BOTH (alternate sentences)
4:40 – 5:00 Outro + thanks	BOTH (smile, slight wave)

Note: Tone reminder: intro and outro are warm and energetic. Problem block is calmer and serious. Demo is curious / discovering tone. Code peek is proud / quietly confident.

Section 4 — The Full 5-Minute Script (word for word)

Total target: ~720 words (~ 4:48 at 150 words per minute). Speakers are marked clearly — follow the order. On-screen cues tell you what should be visible during each line.

0:00 – 0:15 • Hook intro • Both faces on camera

[ON SCREEN] Scene: Both Faces (no screen capture). Title card overlay 'Talk to Data — Theme 1' with both names appears for 3 seconds, then fades out.

FRIEND 1 (Naman) (smile, energetic)

Hi everyone! I'm [Friend 1's full name]...

FRIEND 2 (Your friend) (energetic)

...and I'm [Friend 2's full name].

FRIEND 1 (Naman)

And we're here with our submission for NatWest Code for Purpose, Theme 1 — meet **Talk to Data**.

FRIEND 2 (Your friend)

A tool that turns your spreadsheets and PDFs into instant, trustworthy answers — no SQL, no data team.

0:15 – 0:45 • Problem • Friend 2 leads, slight serious tone

[ON SCREEN] Cut to a simple slide listing 3 pain points: 'Excel limits', 'BI takes weeks', 'Generic AI hallucinates numbers'. Friend 2's webcam tile in bottom-right.

FRIEND 2 (Your friend)

Every day, business teams sit on heaps of data — sales reports, branch CSVs, policy PDFs. But to answer a simple question like 'why did revenue dip last week?' they have to wait for a data team or fight Excel for hours.

FRIEND 2 (Your friend)

Generic AI chatbots feel tempting — but they hallucinate numbers. For a bank, that's a non-starter. So we asked: can we build a tool a non-technical user can trust *in seconds*?

0:45 – 1:15 • Solution overview • Friend 1 leads, confident

[ON SCREEN] Slide cuts to two side-by-side boxes labelled 'Pandas (numbers)' and 'Gemini (words, optional)'. Friend 1's webcam tile in bottom-right.

FRIEND 1 (Naman)

That's where **Talk to Data** comes in. The user just uploads a CSV or PDF, types a question in plain English, and gets a verified answer with the numbers, charts, sources, and assumptions — all in one screen.

FRIEND 1 (Naman)

Here's the key idea: **all numbers are computed locally with pandas**. AI only writes the **summary on top — and only from the verified result**. That means the user can always trust the table, even if AI is off.

1:15 – 1:50 • Tech stack • Friend 2

[ON SCREEN] Tech stack slide: 6 logo tiles (Python, Streamlit, Pandas, Gemini, Postgres+pgvector, Docker). Friend 2 webcam tile bottom-right.

FRIEND 2 (Your friend)

Under the hood, we use **Python and Streamlit** for a beautiful, instant web UI. **Pandas and NumPy** crunch the numbers locally. **pypdf** extracts PDF text.

FRIEND 2 (Your friend)

For optional AI, we use **Google Gemini's free tier** — and we auto-pick the best model the user's API key can access, so demos don't break on hardcoded model IDs. Finally, **Postgres with pgvector** powers semantic search over PDFs — all running together in **Docker** with one command.

1:50 – 2:30 • Demo — CSV upload + first question • Friend 1 driving

[ON SCREEN] Switch to Main Demo scene: full screen of browser at localhost:8501. Both webcams as small PiP tiles in bottom corners.

FRIEND 1 (Naman)

Let's see it live.

[ON SCREEN] Drag and drop data/sample_hackathon.csv onto the upload area.

FRIEND 1 (Naman)

This is the app running locally. I'll upload our sample CSV — branch sales data. Notice — we instantly see the row count, the inferred metric column, and the category dimension. The app figures the schema out for you.

[ON SCREEN] Click in the question box, type 'Why did revenue change last month?', click Generate Insight.

FRIEND 2 (Your friend)

Now let's ask a real business question — 'Why did revenue change last month?'

FRIEND 1 (Naman)

And in under a second — there it is. A change-driver analysis showing which categories pulled revenue up, which pulled it down, with percentages, the source columns, and the assumptions clearly listed.

2:30 – 3:05 • Demo — verified + AI summary • Friend 2 leads, Friend 1 narrates clicks

[ON SCREEN] Scroll down to show the green 'AI explanation (Gemini, grounded)' card. Then scroll back up, change the question to 'North vs South', click Generate Insight.

FRIEND 2 (Your friend)

And below — Gemini gives us a plain-English summary using **only those verified numbers**. So even our manager can read this without knowing what a percentile is.

FRIEND 1 (Naman)

Let's try a comparison — 'North vs South'.

FRIEND 2 (Your friend)

Our query parser detects the 'X vs Y' pattern, runs an entity comparison, shows a bar chart, and the 10%-gap rule flags whether the difference is practically significant.

3:05 – 3:35 • Demo — PDF + AI Q&A; • Friend 2 leads

[ON SCREEN] Click the (x) on the file uploader, drag a sample PDF onto it. Scroll to the AI Q&A; tab, type a question, click Ask AI.

FRIEND 1 (Naman)

Now the PDF side. I'll upload a sample report.

FRIEND 2 (Your friend)

For PDFs, we have three answer modes. Offline keyword search — works even without internet. AI Q&A; — answers strictly from the document text, never outside knowledge. And if Postgres is running, semantic search over chunked embeddings — powered by Gemini's *text-embedding-004*, stored as 768-dimension vectors in pgvector.

[ON SCREEN] Briefly show one AI answer popping up — read the first sentence aloud as proof.

3:35 – 4:10 • Code peek • Friend 1 leads, quietly proud

[ON SCREEN] Cmd+Tab to Cursor. Show the file tree expanded for src/talk_to_data/. Then open analytics.py for 3 seconds. Then open llm_layer.py and highlight the system prompt lines.

FRIEND 1 (Naman)

A quick peek at the code. After our cleanup, the entire app lives in one folder — **src/talk_to_data/**. *app.py* is the Streamlit entry. *views.py* is pure UI. *analytics.py* has every pandas function — like *change_drivers*, which decomposes period-over-period contribution by category.

FRIEND 1 (Naman)

And *llm_layer.py* is where the AI lives — but notice the system prompt: **'use only facts and numbers from the JSON, no outside knowledge'**. That single line is our trust contract.

4:10 – 4:40 • Why this stands out • Both alternate

[ON SCREEN] Cut back to Both Faces scene. No screen capture for this 30 seconds.

FRIEND 2 (Your friend)

What makes this different? **Graceful degradation**. No API key — verified results still work. No internet — PDF line search still works. No database — all CSV analytics still work. Trust is the default; AI is the bonus.

FRIEND 1 (Naman)

Plus we have **transparent UX** — sources, assumptions, and routing confidence are shown *before* any AI text. And every insight is exportable as a JSON audit trail.

4:40 – 5:00 • Outro • BOTH

[ON SCREEN] End-screen lower-third with both names + project name + GitHub URL. Soft music swell up, then fade.

BOTH (*smile, slight wave*)

Thank you for watching!

FRIEND 1 (*Naman*)

I'm **[Friend 1's full name]**...

FRIEND 2 (*Your friend*)

...and I'm **[Friend 2's full name]**.

FRIEND 1 (*Naman*)

This was **Talk to Data** — clarity, trust, and speed by design.

Note: Word count: $\approx$ 720 words across all spoken lines — leaves ~10 seconds of breathing room for pauses, reactions, and music tails. If you go over 5 minutes, the first cut is the second tech-stack sentence; second cut is the second CSV demo question.

Section 5 — Recording Day, Step by Step

5.1 The 60-minute pre-shoot checklist

- Both presenters: water bottle filled, washroom done, throat cleared.
- Phones: Do Not Disturb on, face-down on the desk.
- Laptop: notifications silenced, Slack/WhatsApp closed, Wi-Fi stable.
- OBS: open the 'Main Demo' scene; run 30-sec test, confirm green VU bars on mic.
- Streamlit: running in browser, sample CSV pre-loaded once just to confirm it works, then refresh to clear it.
- Cursor: project open, tabs in the documented order.
- Slides: open in another browser window, presenter view enabled, F5 ready.
- Sample PDF: visible in Finder window so you can drag it instantly.
- OBS hotkey for record: confirmed working (test record / stop).
- Lamp on, hair tidied, shirts smooth — look at yourself in the webcam preview once.

5.2 The recording sequence (do these in this order!)

Take 1 — Slides only screen recording (≈ 30 sec)

- OBS scene: Main Demo, but you'll only show the slides.
- Record yourself flipping from title card → tech stack slide → blank. No audio needed; you'll voice-over later.
- Stop recording. Save with name 'Take_01_Slides.mp4'.

Take 2 — Live demo screen recording (≈ 90 sec)

- OBS scene: Main Demo. Friend 1 drives the mouse and narrates softly so you have a reference timing track.
- Friend 2 stays silent during this take — only screen + Friend 1's voice.
- Hit OBS record. Then in the browser: drag CSV → ask Q1 → scroll to AI answer → ask 'North vs South' → upload PDF → run AI Q&A; → stop recording.
- Don't worry about your narration being perfect; we'll re-record voice in Take 3.
- Save as 'Take_02_Demo.mp4'.

Take 3 — Webcam + voice (full script, ≈ 5 min)

- OBS scene: Both Faces (or Main Demo if you want PiP webcams already composed).
- Both presenters in frame, OBS recording. Read the full script from Section 4 in order.
- If you make a mistake, don't stop. Pause for 2 seconds, restart that line. Editing will fix it.
- Do at least 2 full passes — pick the best later.

- Save as 'Take_03_Webcam_Pass1.mp4', 'Take_03_Webcam_Pass2.mp4', etc.

Take 4 — B-roll / reaction shots (optional, ≈ 30 sec)

- Both presenters smiling, nodding, looking at screen, light laughter — 5-second clips each.
- These are sprinkled in editing during the demo to add life.
- Save as 'Take_04_Broll.mp4'.

5.3 Common issues and quick fixes during recording

- **Echo / hollow audio** → Use earphones with a mic; or move closer to the mic; or hang a folded blanket behind you.
- **Webcam grainy** → Add light directly in front of the face. Bump exposure in OBS source settings.
- **OBS dropping frames** → Lower output to 1280×720 30fps; or change encoder to Apple VT (Mac).
- **Streamlit feels slow** → Restart ``streamlit run ...``; close extra browser tabs.
- **You forget a line** → Pause, take a breath, repeat the previous sentence and continue.

Section 6 — Editing Tutorial (DaVinci Resolve, free)

If DaVinci feels heavy, the same idea works in CapCut for Mac — every step has an equivalent there. We use Resolve in the example because it produces a more polished result and is free.

6.1 First-time setup

1. Download DaVinci Resolve from blackmagicdesign.com (free version).
2. Install. First launch may take a minute.
3. **Project Manager** → New Project → name 'Talk to Data Submission'.
4. **File** → **Project Settings** → Master Settings → Timeline resolution 1920×1080, frame rate 30.

6.2 Import all the takes

1. Open the **Media** tab (bottom strip).
2. Drag your 4 OBS files into the Media Pool.
3. Right-click in Media Pool → New Bin. Create five bins:
 - 01 Webcam — Take_03_Webcam_Pass1, Pass2...
 - 02 Demo — Take_02_Demo
 - 03 Slides — Take_01_Slides
 - 04 BRoll — Take_04_Broll
 - 05 Music — your background music + any SFX

6.3 Build the timeline (the core workflow)

1. Open the **Edit** tab.
2. Drag the best webcam take to **V1** (Video track 1) and **A1** (Audio 1).
3. Right-click the clip in V1 → Decompose Compound Clip if it's grouped — make sure video and audio split.
4. Watch through it. Mark the **start of demo** (where Friend 1 says 'Let's see it live').
5. Drag the demo screen recording onto **V2** (Video 2), aligned at that point.
6. Drag the slides screen recording onto V2 too, placed at 0:15–1:50 area where the script needs slides.

6.4 Picture-in-picture during demo

1. Click the demo clip on V2. Open **Inspector** (right side).

2. Demo clip stays full screen. The webcam below it is hidden during this segment **except** for a small PiP tile.
3. Method A (recommended): cut the webcam clip to take just a 200×200-pixel tile.
4. Method B: keep the webcam clip on V1 full size, then place the demo on V2 with these Inspector → Transform values: *Zoom X 0.78, Zoom Y 0.78, Position X 110, Y -100* — this leaves a strip on the left where the webcam shows through.
5. For a cleaner effect, use a fresh PiP layer: copy webcam clip up to V3, scale to Zoom 0.22, Position X 720, Y -380. Add a 4-pixel white border via Effects → ResolveFX Stylize → Stroke.

6.5 Cut out filler and silence

- Press **B** to enter Blade mode. Click on every 'um', long pause, or repeated take to slice.
- Press **A** to return to Selection mode. Select the bad slice → press Backspace.
- Hold **Shift+Backspace** to ripple-delete (closes the gap).
- Aim for snappy cuts; never leave more than 0.7 sec of silence.

6.6 Crossfades and transitions

1. Between major blocks (intro → problem, problem → solution, etc.) add a **Cross Dissolve** of 6–10 frames.
2. From the Effects library, drag **Cross Dissolve** onto each cut.
3. Don't overuse fancy transitions — they look amateur. 90% of cuts should be straight cuts.

6.7 Audio polish

1. Click the audio clip → Inspector → **Equalizer**: enable High-pass at 80 Hz to kill rumble.
2. Add **Compressor**: ratio 3:1, threshold around -18 dB.
3. Add **Voice Isolation** (Studio version) or use 'Noise Reduction' in the Fairlight tab if there's hiss.
4. Drag background music onto **A3**. Lower its volume to about **-25 dB** in Inspector.
5. Right-click the music clip → Smart Filter → 'Side-Chain to A1' so music ducks under your voice automatically.
6. Add 1-second fade-in at the start, 2-second fade-out at the end of the music.

6.8 Lower thirds (names)

1. Effects library → Titles → drag **Lower Third** onto track V4 above everything.
2. Edit text in Inspector: 'Friend 1's Name — Project Lead' (or 'Co-Presenter').
3. Show for 3–5 seconds the first time each presenter speaks (around 0:03 and 0:18).
4. Use a consistent font (e.g. Inter, Helvetica Neue) and the project's blue (#0ea5e9) accent.

6.9 Title card and end card

1. Open Fusion tab (advanced) OR simply use Effects → Titles → 'Text+'.
2. Title card (first 3 seconds): 'Talk to Data — Theme 1' big, both names below in smaller font.
3. End card (last 5 seconds): 'Thank you' + project name + GitHub URL.
4. Background: solid #0f172a navy with a subtle radial gradient toward #1e3a8a.

6.10 Captions / subtitles (optional but strongly recommended)

1. Right-click on the timeline ruler → **Create Subtitles from Audio**.
2. Resolve transcribes — review every line and fix any wrong words.
3. Style: white text, black outline 2 px, font size 36, position bottom-third.
4. Captions massively help non-native English viewers and people who watch on mute.

6.11 Color correction (quick pass)

1. Open the **Color** tab.
2. Select all webcam clips. Use the Color Wheels: lift Lift slightly toward warm, push Gain slightly toward cool — keeps skin tones natural.
3. Add a small Saturation boost (about +5).
4. If clips were handheld, apply Stabiliser (Color tab → Stabilization → Smooth).

6.12 Final QC inside the editor

- Watch the entire video without pausing. Note timestamps of any issue.
- Fix issues. Watch again.
- Length must be under 5:00. Sweet spot: 4:45–4:55.
- Listen with earphones — confirm no clipping (no red on the audio meter).
- Confirm both presenters' faces are well-lit and audio is balanced between speakers.

6.13 Export the final file

1. Open the **Deliver** tab.
2. Preset: **YouTube 1080p**.
3. Format: MP4. Codec: H.264. Audio: AAC 192 kbps, 48 kHz.
4. Quality: Restrict to ~12,000 kbps.
5. Filename: **Talk_to_Data_Submission_Final.mp4**.
6. Add to Render Queue → Start Render. Wait until 100%.

Section 7 — Final Quality Gate

Open the exported file and tick every box before submitting.

- ☐ Audio: clear, no clipping (no red meter), music quietly behind the voice.
- ☐ Video: stable, colour consistent, both faces visible during intro/outro.
- ☐ Text: zero spelling mistakes in titles, lower thirds, and captions.
- ☐ Length: between 4:30 and 5:00. Aim for 4:48–4:55.
- ☐ Both names visible at intro, outro, and at least one lower third in the middle.
- ☐ Demo screen text is readable (zoom in on a phone to test — if you can read it, judges can).
- ☐ No accidental personal info on screen — no other browser tabs, no notification popups, no exposed Gemini key.
- ☐ File size under 500 MB ideally — fits most upload limits comfortably.
- ☐ Watched once with earphones, once with laptop speakers — both sound fine.

Section 8 — Submission

Most hackathons want a YouTube/Drive link. If they want a direct upload, the MP4 you exported is ready. Steps if it's a YouTube link:

1. Sign in to youtube.com → Create → Upload Video.
2. Title: 'Talk to Data — NatWest Code for Purpose Theme 1 Submission'.
3. Description: 1-line summary, names, GitHub link.
4. Visibility: **Unlisted** (anyone with the link can watch, but it doesn't appear in search).
5. After upload completes, copy the link, open it in **incognito** to confirm it works without your account.
6. Submit the link via the hackathon portal.

Note: Keep your final MP4 + project ZIP in two places (your laptop + Google Drive) so you can re-upload if anything fails.

Section 9 — Backup Plans (when something goes wrong)

If the video runs over 5 minutes

- Cut the second tech-stack sentence (the one starting 'For optional AI...').
- Cut the second CSV demo example ('North vs South' segment).
- Tighten the outro by removing the line repetition of names.

If the video runs under 4:20

- Add a 10-second 'What's next' slide before the outro: scaling, audit logs, more connectors.
- Slow down the demo a touch — let one verified result sit on screen for 3 extra seconds.

If the audio quality is bad

- Re-record only the audio against the existing video (audio replacement). It's quicker than re-shooting.
- In Resolve, mute the original audio track, drag the new audio onto A2, sync visually with mouth movement.

If a demo step glitches mid-recording

- Take a screenshot of the same screen later, place it on V2 for those few seconds, voice over it.

If you both look stiff

- Re-record the intro and outro standing up — energy travels through your body to your voice.
- Have a friend off-camera making faces; both presenters react to them just before recording starts.

Section 10 — Tips That Make the Video Engaging

- Smile during intro and outro. People mirror what they see.
- Use hand gestures within frame — keeps the eye moving without moving the camera.
- Vary pitch on key words: 'Trust', 'verified', 'instant', 'first', 'only'.
- Look at the camera lens, not the screen, when you're speaking to the audience.
- Pause briefly (0.5 s) after a key claim — gives it weight.
- React on camera to your co-host: nods, brief 'yeah', light smile.
- Cut to demo footage early — judges trust what they can see working.
- Keep on-screen captions minimal: 4–6 words max per line.
- End with a clear value line, not a thank-you alone — 'Talk to Data — clarity, trust, speed.'
- Don't apologise on camera. Don't say 'this might not work'. Confidence sells.

Appendix A — Exactly Which Code to Show on Screen

When Friend 1 switches from the demo to the editor, here are the exact files and the exact lines to highlight. Keep each visible on screen for ~3–4 seconds — long enough to read but not long enough to read every line.

A1. Folder tree (open the Cursor sidebar)

- Show `src/talk_to_data/` expanded so all 12 files are visible.
- Mention: 'one folder, one entry, one place to look'.

A2. `analytics.py` — the `change_drivers` function

- Open `src/talk_to_data/analytics.py`. Scroll to `def change_drivers`.
- Highlight the line that computes `contribution_pct` — this is your 'why did revenue change' answer.

A3. `llm_layer.py` — the trust contract

- Open `src/talk_to_data/llm_layer.py`. Scroll to `CSV_ENRICH_SYSTEM`.
- Highlight: *'Use ONLY facts and numbers from the JSON. No outside knowledge.'* — this is the line you read aloud.

A4. `pgvector_store.py` — the schema (optional, only if time permits)

- Show `_DDL_TABLE` — the `pdf_chunks` table with `vector(768)` column.
- One sentence: 'Same Gemini embedding stored in Postgres for semantic search.'

Appendix B — Word-count budget per block

Use this if you want to time-trim or expand any block.

Hook intro (0:00–0:15)	≈ 35 words
Problem (0:15–0:45)	≈ 75 words
Solution overview (0:45–1:15)	≈ 70 words
Tech stack (1:15–1:50)	≈ 80 words
Demo + commentary (1:50–3:35)	≈ 240 words
Code peek (3:35–4:10)	≈ 90 words
Why it stands out (4:10–4:40)	≈ 70 words
Outro (4:40–5:00)	≈ 35 words
TOTAL target	≈ 695–720 words (≈ 4:48 spoken)

Appendix C — Day-of cheat sheet (print this!)

BEFORE you press record

- DND on, all notifications closed, browser fullscreen.
- Streamlit running, sample CSV + sample PDF in Finder.
- Cursor tabs ordered: README, app, views, analytics, llm_layer, pgvector_store.
- Slides ready: title card + tech stack.
- OBS scenes set: 'Both Faces' and 'Main Demo'. Hotkey Cmd+R configured.
- Mic peaks at -12 to -6 dB. Lamp on. Water near you.

DURING recording

- Take 1: Slides only — record the slide flips.
- Take 2: Demo — Friend 1 narrates clicks (rough audio).
- Take 3: Webcam — full script, both faces, 2 passes.
- Take 4 (optional): B-roll — smiles, nods, reactions.

AFTER recording, in DaVinci

- Import 4 takes into bins. Build timeline V1=webcam, V2=screen, V3=PiP, V4=titles, A3=music.
- Cut filler. Add cross dissolves at major joins. Add lower thirds for names.
- Audio: HP filter, compressor, music ducked to -25 dB with 1s fade-in / 2s fade-out.
- Captions: auto-transcribe, fix words, white text + black outline.
- Title card 3 sec, end card 5 sec. Watch end-to-end. Length 4:48–4:55.
- Deliver: MP4 H.264 1080p, AAC 192k, ~12 Mbps. Filename Talk_to_Data_Submission_Final.mp4.

FINAL

- Watch once with earphones, once on speaker. Re-export if needed.
- Upload to YouTube as Unlisted. Test link in incognito. Submit.
- Backup the MP4 + project ZIP to Drive.

Note: You will not get this perfect on the first take. That's normal. The script + this cheat sheet is what makes the second and third takes look easy. Trust the process — and have fun.